

# Database Technology

## Topic 9: Concurrency Control

Olaf Hartig

[olaf.hartig@liu.se](mailto:olaf.hartig@liu.se)

# Goal

- Preserve **Isolation** of the ACID properties

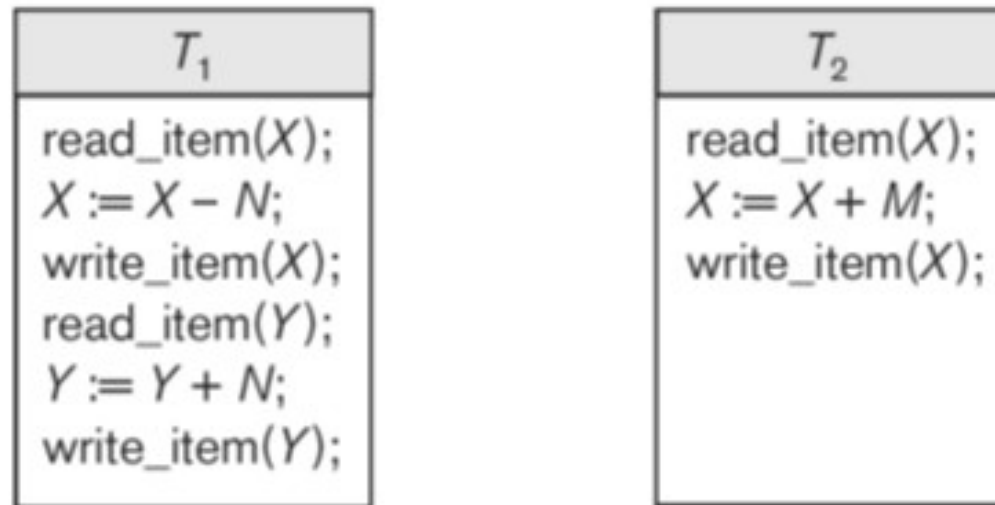


# Notation and Initial Concepts

# Transaction Notation

- Focus on read and write operations
  - For instance,  $w_5(Z)$  means that transaction 5 writes data item  $Z$
- $b_i$  and  $e_i$  specify transaction boundaries (begin and end)
  - $i$  specifies a unique transaction identifier (TID)

- Example:



- $T_1$ :  $b_1, r_1(X), w_1(X), r_1(Y), w_1(Y), e_1$
- $T_2$ :  $b_2, r_2(X), w_2(X), e_2$

# Schedule

- Sequence of interleaved operations from multiple TAs

- Example:

	at ATM window #1	at ATM window #2
1	read_item(savings);	
2	savings = savings - \$100;	
3		read_item(checking);
4	write_item(savings);	
5	read_item(checking);	
6		checking = checking - \$20;
7		write_item(checking);
8	checking = checking + \$100;	
9	write_item(checking);	
10		dispense \$20 to customer;

–  $S: b_1, r_1(s), b_2, r_2(c), w_1(s), r_1(c), w_2(c), w_1(c), e_1, e_2$

# Quiz

What can be concluded from the following schedule?

...,  $r_3(\text{EMPLOYEE})$ ,  $b_4$ ,  $w_2(\text{STUDENT})$ , ...

- A: Some employee has read a student record.
- B: A transaction has read some data and then written it back.
- C: At least three transactions were running concurrently.
- D: All of the above.
- E: None of the above.

# Serial Schedules

- Definition: a schedule is *serial* if the operations of any TA are executed directly one after the other
  - i.e., no interleaving of operations from different TAs
- Characteristics:
  - Serial schedules trivially guarantee the isolation property
  - For  $n$  transactions, there are  $n!$  serial schedules
  - Each of them produces a correct result (assuming the consistency preservation property)
  - However, not all of them might produce the *same* result
    - For instance, If two people try to reserve the last seat on a plane, only one gets it. The serial order determines which one. The two orderings have different results, but either one is correct.

# Serial Schedules (cont'd)

Serial schedules are *not feasible* for performance reasons:

- Long transactions force other transactions to wait
- When a transaction is waiting for disk I/O or any other event, system cannot switch to other transaction
- Solution: allow *some* interleaving  
(without sacrificing correctness!)



# Acceptable Interleavings

(Serializability)

# Conflicts

- Executing some operations in a different order causes a different outcome
  - ...  $r_1(X)$ ,  $w_2(X)$ , ... vs. ...  $w_2(X)$ ,  $r_1(X)$ , ...  
 $T_1$  will read a different value for  $X$
  - ...  $w_1(Y)$ ,  $w_2(Y)$ , ... vs. ...  $w_2(Y)$ ,  $w_1(Y)$ , ...  
value for  $Y$  after both operations will be different
- Note that two read operations do not have this issue
  - ...  $r_1(Z)$ ,  $r_2(Z)$ , ... vs. ...  $r_2(Z)$ ,  $r_1(Z)$ , ...  
both TAs read the same value of  $Z$

# Conflicts and Equivalence

*Definition:* Two operations **conflict** if

1. they access the same data item  $X$ ,
2. they are from two different transactions, and
3. at least one of them is a write operation.

*Definition:* Two schedules are **conflict equivalent** if the relative order of *any two conflicting operations* is the same in both schedules.

# Serializability

*Definition:* A schedule with  $n$  transactions is **serializable** if it is conflict equivalent to *some* serial schedule of the same  $n$  transactions.

- Serializable schedule “correct” because equivalent to some serial schedule, and any serial schedule acceptable
  - Transactions see data as if they were executed serially
  - Transactions leave DB state as if they were executed serially (hence, serializable schedules will leave the database in a consistent state)
- Efficiency achievable through interleaving and concurrent execution

# Testing Serializability


- Construct a *serialization graph* for the schedule
  - Node for each transaction in the schedule
  - Direct edge from  $T_i$  to  $T_j$  if some read or write operation in  $T_i$  appears before a conflicting operation in  $T_j$
- A schedule is serializable if and only if its serialization graph has not cycles

# Example

- Consider the following schedule

$S: b_1, r_1(X), b_2, r_2(Y), w_1(X), b_3, w_2(Y), e_2, r_1(Y), r_3(X), e_3, w_1(Y), e_1$

- Serialization graph of  $S$ :



- No cycles! Hence,  $S$  is serializable.
  - Equivalent to the following serial schedule:

$S': b_2, r_2(Y), w_2(Y), e_2, b_1, r_1(X), w_1(X), r_1(Y), w_1(Y), e_1, b_3, r_3(X), e_3$

The serial schedule  $S'$  is shown with dashed blue lines indicating the execution order of transactions. The sequence of operations is:  $b_2, r_2(Y), w_2(Y), e_2$  (labeled  $T_2$ ), followed by  $b_1, r_1(X), w_1(X), r_1(Y), w_1(Y), e_1$  (labeled  $T_1$ ), followed by  $b_3, r_3(X), e_3$  (labeled  $T_3$ )).

# Quiz

## Remember

- If the initial value of checking is \$500, what value does it have after the following interleaved execution completes?

	at ATM window #1	at ATM window #2
1	read_item(savings);	
2	savings = savings - \$100;	
3		read_item(checking);
4	write_item(savings);	
5	read_item(checking);	
6		checking = checking - \$20;
7		write_item(checking);
8	checking = checking + \$100;	
9	write_item(checking);	
10		dispense \$20 to customer;

**A:** \$480

**B:** \$500

**C:** \$580

**D:** \$600

– S:  $b_1, r_1(s), b_2, r_2(c), w_1(s), r_1(c), w_2(c), w_1(c), e_1, e_2$

# Key Question

Can we make sure that we only get serializable schedules?



# Locking Techniques for Concurrency Control


# Database Locks

- Locks can be used to ensure that conflicting operations cannot occur
- **Exclusive lock** for writing, **shared lock** for reading
  - Transaction cannot read item without first getting a shared or an exclusive lock on it
  - Transaction cannot write item without first getting exclusive lock on it



# Database Locks (cont'd)

- Request for lock may cause transaction to **block** (wait) because write lock is exclusive
  - Any lock on  $X$  (read or write) cannot be granted if some other transaction holds write lock on  $X$
  - Write lock on  $X$  cannot be granted if some other transaction holds *any* lock on  $X$




- Blocked transactions are unblocked and granted the requested lock when conflicting transaction(s) release their lock(s)



# Two-Phase Locking (2PL)

*Definition:* A transaction follows the two-phase locking (2PL) protocol if *all* of its `read_lock()` and `write_lock()` operations come before its first `unlock()` operation.



- A transaction that follows the 2PL protocol has an *expansion phase* and a *shrinking phase*
- If all transactions in a schedule follow the 2PL protocol, then the schedule is serializable



# Deadlock

- Two or more transactions wait for one another to unlock some data item
  - $T_i$  waits for  $T_j$  waits for ... waits for  $T_n$  waits for  $T_i$
- Deadlock prevention:
  - Conservative 2PL protocol: Wait until you can lock all the data to be used beforehand
  - Wait-die
  - Wound-wait
  - No waiting
  - Cautious waiting
- Deadlock detection:
  - Wait-for graph
  - timeouts

# Starvation

- A transaction is not executed for an indefinite period of time while other transactions are executed normally
  - e.g.,  $T$  waits for write lock and other TAs repeatedly grab read locks before all read locks are released
- Starvation prevention:
  - First-come-first-served waiting scheme
  - Wait-die
  - Wound-wait
  - etc.

# Summary

# Summary

- Characterizing schedules based on serializability
  - Serial and non-serial schedules
  - Conflict equivalence of schedules
  - Serialization graph
- Two-phase locking
  - Guarantees conflict serializability
  - Possible problems: deadlocks and starvation



[www.liu.se](http://www.liu.se)