

Database Technology

Topic 10: Database Recovery

Olaf Hartig

olaf.hartig@liu.se

Types of Failures

Database may become unavailable for use due to:

- Transaction failures
 - e.g., incorrect input, deadlock, incorrect synchronization
 - Result: transaction abort
- System failures
 - e.g., application error, operating system fault
- Media failures
 - e.g., RAM failure, disk head crash, power disruption


Focus of the rest of the lecture:

- We wish to recover from system failures
 - Recovery from media failures similar, but may need to restore database files from *backup*

Situation after System Failure

- DBMS is halted abruptly
- Processing of in-progress SQL commands halted abruptly
- Connections to application programs (clients) are broken
- States of executing programs unknown
- Contents of memory buffers are lost
- Database files are *not* damaged

Problem Situation Example



- T_1 , T_2 , and T_3 have committed
- T_4 and T_5 still in progress
- Any of the transactions might have written data
- Some (unknown) subset of the writes have been flushed to disk

Purpose of Database Recovery

- Bring the database into the most recent consistent state that existed prior to a failure
- **Atomicity** and **Durability** of the ACID properties
 - Abort (and restart) TAs active at time of failure
 - Ensure changes made by committed TAs are not lost
- Complication due to database execution model:
 - Data items packed into I/O blocks (pages)
 - At time of write updated data first stored in main memory buffer
 - Actually written to disk some time later

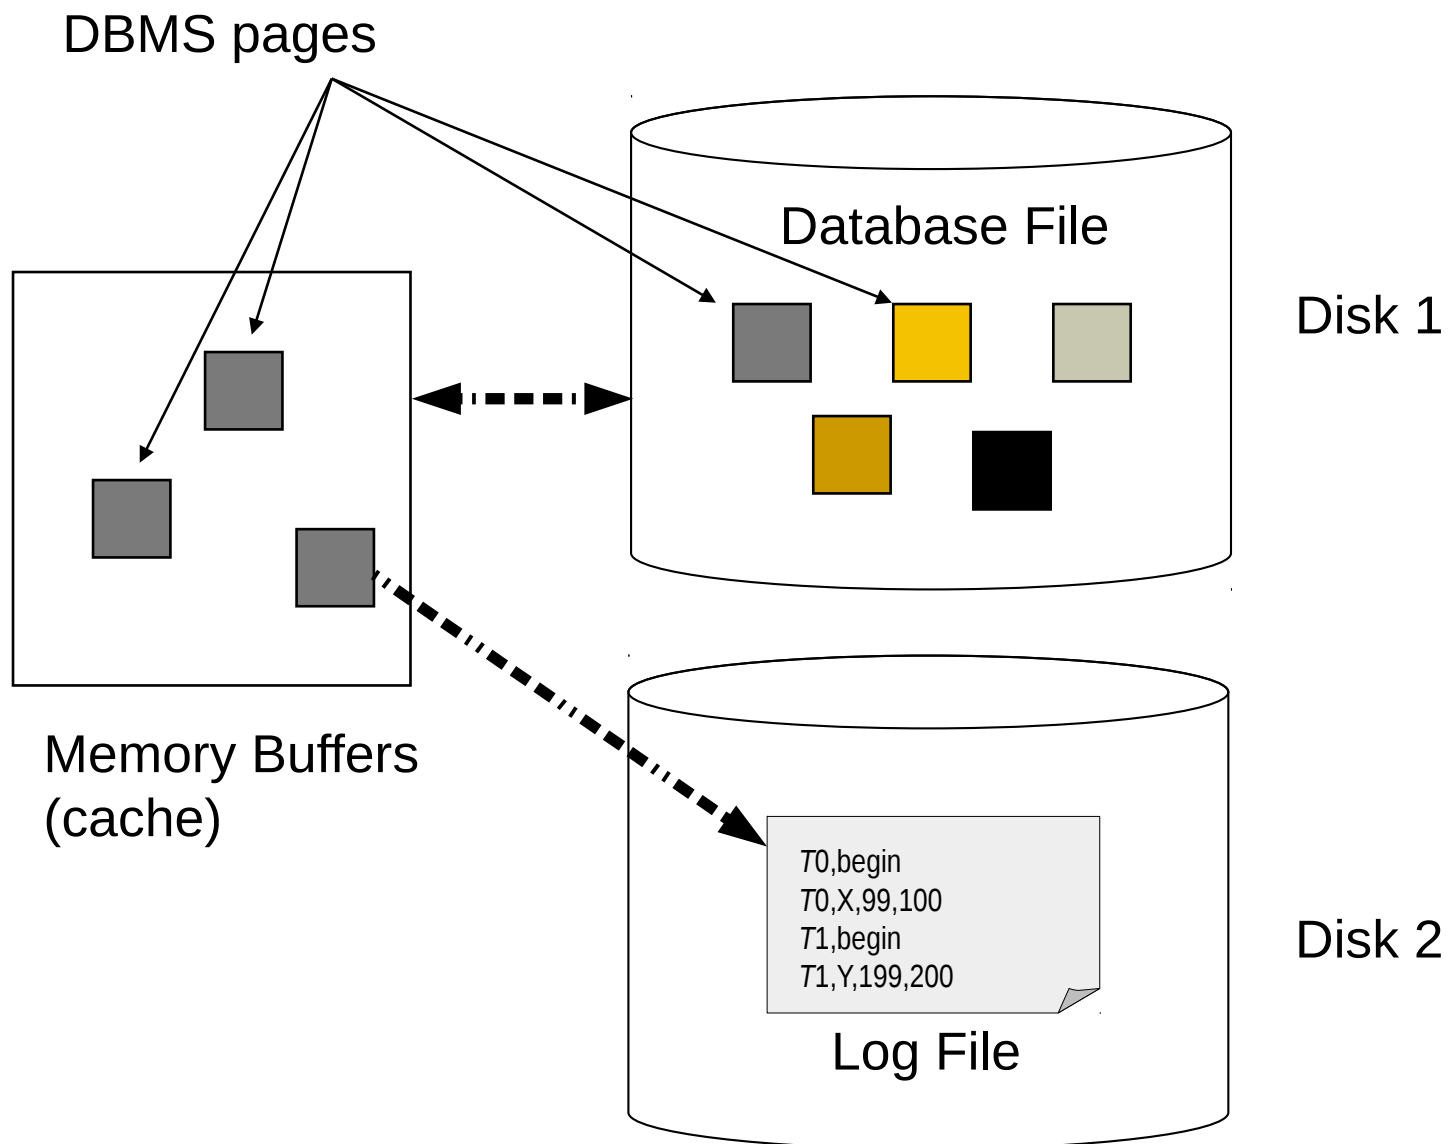
Logging

System Log

- Append-only file
 - Keep track of all operations of all transactions
 - In the order in which operations occurred
- Stored on disk
 - Persistent except for disk or catastrophic failures
 - Periodically backed up (to guard against disk and catastrophic failures)
- Log buffer in main memory
 - Holds log records that are appended to the log
 - Occasionally whole buffer appended to end of log file on disk (flush)



Storage Structure



Log Records

- [**start_transaction**, T]
 - Transaction T has started execution
- [**write_item**, T , X , old_value , new_value]
 - Transaction T has changed the value of item X from old_value to new_value
 - old_value (before image) needed to **undo**(X)
 - new_value (after image) needed to **redo**(X)
- [**commit**, T]
 - T has completed successfully and committed
 - Effects (writes) of T must be durable
- [**abort**, T]
 - T has been aborted
 - Effects (writes) of T must be ignored and undone

Commit Point

- A transaction reaches its commit point when:
 1. all of its operations are executed, and
 2. all its log records are flushed to disk
(where the last is the commit record)
- Beyond its commit point
 - the transaction is said to be *committed*, and
 - its effect must be permanently recorded in the DB



Write-Ahead Logging (WAL)

- Used to ensure that the log
 - is consistent with the DB, and
 - can be used to recover the DB to a consistent state
- Two rules:
 1. Log record(s) for a page must be written before corresponding page is flushed to disk, and
 2. All log records must be written before commit
- Rule 1 for atomicity
 - each operation is known and can be undone if needed
- Rule 2 for durability
 - the effect of a committed transaction is known

Recovery Process

Recovery with Deferred Update

- Updating the DB on disk after each change is inefficient
- **Deferred update:**
 - Updates of a transaction T are written to disk *after (but not necessarily immediately after)* T has reached commit point
- No need to undo changes of non-committed transactions
- Need to redo the changes of committed transactions
- **NO-UNDO/REDO recovery algorithm:**
 - Create a list of active (i.e., non-committed) transactions and a list of committed transactions
 - REDO all the write-item operations of all the TAs in the second list in the order in which they appear in the log (use *after image* from the log records)

Example

NO-UNDO

REDO: T1, T4

start-transaction T1
write-item T1, D, 10, 20
commit T1

start-transaction T4
write-item T4, B, 10, 20
write-item T4, A, 5, 10
commit T4

start-transaction T2


~~write-item T2, B, 20, 15~~ - ignore

start-transaction T3

~~write-item T3, A, 10, 30~~ - ignore


~~write-item T2, D, 20, 25~~ - ignore

CRASH



Checkpointing

- To save redo effort, use **checkpoints**
- Occasionally flush data buffers using the five steps:
 1. Suspend execution of transactions temporarily
 2. Force-write modified buffer data to disk
 3. Append [checkpoint] record to log
 4. Flush log to disk
 5. Resume normal transaction execution
- During recovery, redo required only for log records appearing after [checkpoint] record



Example with Checkpoint

NO-UNDO

REDO: ~~T1~~, T4

start-transaction T1
~~write-item T1, D, 10, 20~~ *ignore*
commit T1
checkpoint
start-transaction T4
write-item T4, B, 10, 20
write-item T4, A, 5, 10
commit T4
start-transaction T2
~~write-item T2, B, 20, 15~~ *ignore*
start-transaction T3
~~write-item T3, A, 10, 30~~ *ignore*
~~write-item T2, D, 20, 25~~ *ignore*
CRASH

T1

T4

T2

T3

checkpoint

crash

Recovery with Immediate Update 1

- **Immediate update:**
 - Updates of a transaction may be written to disk *before* the transaction commits (with the log records for such updates being written out first, i.e., write-ahead logging)
- Additional requirement: all updates of a transaction T must be written to disk before the commit point of T
 - No need to redo changes of committed transactions
 - Need to undo changes of non-committed transactions
- **UNDO/NO-REDO recovery algorithm:**
 - Create a list of active (i.e., non-committed) transactions
 - UNDO all the write-item operations of all the TAs in the list in the *reverse* order in which they appear in the log (use *before image* from the log records)

Example

UNDO: T2, T3

NO REDO

start-transaction T1

~~write-item T1, D, 10, 20~~ *ignore*

commit T1

~~checkpoint~~ *not needed*

start-transaction T4

~~write-item T4, B, 10, 20~~ *ignore*

~~write-item T4, A, 5, 10~~ *ignore*

commit T4

start-transaction T2


write-item T2, B, 20, 15

start-transaction T3

write-item T3, A, 10, 30

write-item T2, D, 20, 25

CRASH



Recovery with Immediate Update 2


- No additional requirements
- Then:
 - Need to redo changes of committed transactions
 - Need to undo changes of non-committed transactions
- **UNDO/REDO recovery algorithm:**
 - Create a list of active (i.e., non-committed) transactions and a list of committed transactions *since last checkpoint*
 - UNDO all the write-item operations of all the TAs in the first list in the *reverse* order in which they appear in the log (use *before image* from the log records)
 - REDO all the write-item operations of all the TAs in the second list in the order in which they appear in the log (use *after image* from the log records)

Example

UNDO: T2, T3

REDO: T4

start-transaction T1
~~write-item T1, D, 10, 20~~ *ignore*
commit T1
checkpoint
start-transaction T4
write-item T4, B, 10, 20
write-item T4, A, 5, 10
commit T4
start-transaction T2
write-item T2, B, 20, 15
start-transaction T3
write-item T3, A, 10, 30
write-item T2, D, 20, 25
CRASH



Summary

Summary

- Transaction log
- Transaction roll-back (undo) and roll-forward (redo)
- Checkpointing

www.liu.se